

# Lista II: Q1

Alberson Miranda

2022-11-27

```
# configurações
knitr::opts_chunk$set(
  fig.output = "70%"
)

# reproducibilidade
set.seed(1)

# pacotes
pacman::p_load(
  "ggplot2",
  "tsibble",
  "fable",
  "feasts",
  "fabletools",
  "urca"
)
```

## 1 MODELAGEM BOX-JENKINS: SÉRIE I

O primeiro passo é a importação e visualização da série. Como não há informação sobre o período, usarei diário e tentarei identificar a partir de um padrão sazonal, se houver.

```
# importando dados
load("data/lista II.RData")
data = data.frame(
  value = conjunto1[, 1],
  index = seq(
    as.Date("2000-01-01"),
```

```

    by = 1,
    length.out = length(conjunto1[, 1])
  )
) |> tsibble(index = index)


```

A série é compacta, ou seja, de amplitude baixa, não requerindo transformação para redução de variância.

```

# plot série
autoplot(data, .vars = value)

```



O segundo passo é testar se a série é estacionária no primeiro momento. Não há evidências de raiz unitária tanto nos testes quanto nos gráficos de autocorrelação. Para os testes ADF, iniciei com a especificação com tendência. Não sendo significativo o coeficiente  $\tau_t$ , passei para a especificação com *drift*, sendo tanto o intercepto quanto  $z.lag.1$  significativos, adoto esta como a correta especificação e, assim como nos testes de Phillips-Perron e KPSS, não há indicativo de raiz unitária.

```

# KPSS test
data |>
  features(value, unitroot_kpss)

```

```
# A tibble: 1 x 2
  kpss_stat kpss_pvalue
    <dbl>      <dbl>
1    0.108      0.1
```

```
# Phillips-Perron test
data |>
  features(value, unitroot_pp)
```

```
# A tibble: 1 x 2
  pp_stat pp_pvalue
    <dbl>      <dbl>
1   -6.43      0.01
```

```
# Augmented-Dickey-Fuller test
data |>
  (\(x) ur.df(x$value, selectlags = "AIC", type = "trend", lags = 12))() |>
  summary()
```

```
#####
# Augmented Dickey-Fuller Test Unit Root Test #
#####
```

Test regression trend

Call:  
lm(formula = z.diff ~ z.lag.1 + 1 + tt + z.diff.lag)

Residuals:

Min	1Q	Median	3Q	Max
-0.0102211	-0.0022799	0.0000216	0.0020829	0.0102631

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	1.141e+01	2.057e+00	5.549	9.87e-08 ***
z.lag.1	-3.424e-01	6.171e-02	-5.549	9.87e-08 ***
tt	3.313e-06	4.807e-06	0.689	0.492
z.diff.lag	-8.474e-03	7.397e-02	-0.115	0.909

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.00356 on 184 degrees of freedom

Multiple R-squared: 0.171, Adjusted R-squared: 0.1574

F-statistic: 12.65 on 3 and 184 DF, p-value: 1.487e-07

Value of test-statistic is: -5.5495 10.2859 15.411

Critical values for test statistics:

	1pct	5pct	10pct
tau3	-3.99	-3.43	-3.13
phi2	6.22	4.75	4.07
phi3	8.43	6.49	5.47

```
data |>
  (\(x) ur.df(x$value, selectlags = "AIC", type = "drift", lags = 12))() |>
  summary()
```

```
#####
# Augmented Dickey-Fuller Test Unit Root Test #
#####
```

Test regression drift

Call:

```
lm(formula = z.diff ~ z.lag.1 + 1 + z.diff.lag)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.0101264	-0.0022878	-0.0000252	0.0020877	0.0103523

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	11.27905	2.04454	5.517	1.15e-07 ***
z.lag.1	-0.33838	0.06134	-5.517	1.15e-07 ***
z.diff.lag	-0.01060	0.07380	-0.144	0.886

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.003555 on 185 degrees of freedom  
Multiple R-squared: 0.1688, Adjusted R-squared: 0.1598  
F-statistic: 18.79 on 2 and 185 DF, p-value: 3.73e-08

Value of test-statistic is: -5.5166 15.2347


Critical values for test statistics:

	1pct	5pct	10pct
tau2	-3.46	-2.88	-2.57
phi1	6.52	4.63	3.81

A seguir, pode-se perceber decaimento na ACF e um pico na PACF, sugerindo um processo AR(1).


```
# ACF  
data |> ACF() |> autoplot()
```

Response variable not specified, automatically selected `var = value`



```
# ACF  
data |> PACF() |> autoplot()
```

Response variable not specified, automatically selected `var = value`



Além do AR(1), também realizei uma *grid search*, que consiste na estimação de todas as combinações possíveis de modelos ARIMA dada uma restrição de coeficientes. Neste caso, como a análise do correlograma sugere um AR(1), optei por uma restrição parcimoniosa, com no máximo 3 coeficientes (AR ou MA) e sem testar modelos integrados, uma vez que foi constatada a estacionaridade. Dentre os modelos estimados, o de menor critério de informação foi o AR(1), no mesmo sentido da análise visual do correlograma.

```
data_fit = data |>
  model(
    ar1 = ARIMA(
      value ~ 1 + pdq(1, 0, 0)
    ),
    search = ARIMA(
      value,
      stepwise = FALSE,
      trace = TRUE,
      order_constraint = p + q + P + Q <= 3 & (constant + d + D <= 1))
  )
```

ARIMA(0,0,0)(0,0,0)[7]+c      -1589.386363

ARIMA(1,0,0)(0,0,0)[7]+c	-1700.363424
ARIMA(2,0,0)(0,0,0)[7]+c	-1698.315341
ARIMA(3,0,0)(0,0,0)[7]+c	-1695.712579
ARIMA(0,0,1)(0,0,0)[7]+c	-1668.586911
ARIMA(1,0,1)(0,0,0)[7]+c	-1698.303988
ARIMA(2,0,1)(0,0,0)[7]+c	-1696.822220
ARIMA(0,0,2)(0,0,0)[7]+c	-1682.955592
ARIMA(1,0,2)(0,0,0)[7]+c	-1696.222233
ARIMA(0,0,3)(0,0,0)[7]+c	-1695.230046
ARIMA(0,0,0)(1,0,0)[7]+c	-1585.461121
ARIMA(1,0,0)(1,0,0)[7]+c	-1696.000612
ARIMA(2,0,0)(1,0,0)[7]+c	-1693.062626
ARIMA(0,0,1)(1,0,0)[7]+c	-1663.280799
ARIMA(1,0,1)(1,0,0)[7]+c	-1693.950922
ARIMA(0,0,2)(1,0,0)[7]+c	-1678.426909
ARIMA(0,0,0)(2,0,0)[7]+c	-1578.757729
ARIMA(1,0,0)(2,0,0)[7]+c	-1688.504270
ARIMA(0,0,1)(2,0,0)[7]+c	-1656.733042
ARIMA(0,0,0)(0,0,1)[7]+c	-1588.036672
ARIMA(1,0,0)(0,0,1)[7]+c	-1699.009375
ARIMA(2,0,0)(0,0,1)[7]+c	-1697.025695
ARIMA(0,0,1)(0,0,1)[7]+c	-1666.716920
ARIMA(1,0,1)(0,0,1)[7]+c	-1696.968352
ARIMA(0,0,2)(0,0,1)[7]+c	-1681.620771
ARIMA(0,0,0)(1,0,1)[7]+c	-1595.850518
ARIMA(1,0,0)(1,0,1)[7]+c	-1697.251094
ARIMA(0,0,1)(1,0,1)[7]+c	-1667.413004
ARIMA(0,0,0)(2,0,1)[7]+c	Inf
ARIMA(0,0,0)(0,0,2)[7]+c	-1586.749289
ARIMA(1,0,0)(0,0,2)[7]+c	-1696.969600
ARIMA(0,0,1)(0,0,2)[7]+c	-1665.570336
ARIMA(0,0,0)(1,0,2)[7]+c	-1593.766397
ARIMA(0,0,0)(0,0,0)[7]	1982.061534
ARIMA(1,0,0)(0,0,0)[7]	Inf
ARIMA(2,0,0)(0,0,0)[7]	Inf
ARIMA(3,0,0)(0,0,0)[7]	Inf
ARIMA(0,0,1)(0,0,0)[7]	1725.595542
ARIMA(1,0,1)(0,0,0)[7]	Inf
ARIMA(2,0,1)(0,0,0)[7]	Inf
ARIMA(0,0,2)(0,0,0)[7]	1529.598405
ARIMA(1,0,2)(0,0,0)[7]	Inf
ARIMA(0,0,3)(0,0,0)[7]	1393.372344
ARIMA(0,0,0)(1,0,0)[7]	Inf

ARIMA(1,0,0)(1,0,0)[7]	Inf
ARIMA(2,0,0)(1,0,0)[7]	Inf
ARIMA(0,0,1)(1,0,0)[7]	Inf
ARIMA(1,0,1)(1,0,0)[7]	Inf
ARIMA(0,0,2)(1,0,0)[7]	Inf
ARIMA(0,0,0)(2,0,0)[7]	Inf
ARIMA(1,0,0)(2,0,0)[7]	Inf
ARIMA(0,0,1)(2,0,0)[7]	Inf
ARIMA(0,0,0)(0,0,1)[7]	1758.879724
ARIMA(1,0,0)(0,0,1)[7]	Inf
ARIMA(2,0,0)(0,0,1)[7]	Inf
ARIMA(0,0,1)(0,0,1)[7]	1555.857332
ARIMA(1,0,1)(0,0,1)[7]	Inf
ARIMA(0,0,2)(0,0,1)[7]	1426.549252
ARIMA(0,0,0)(1,0,1)[7]	Inf
ARIMA(1,0,0)(1,0,1)[7]	Inf
ARIMA(0,0,1)(1,0,1)[7]	Inf
ARIMA(0,0,0)(2,0,1)[7]	Inf
ARIMA(0,0,0)(0,0,2)[7]	1618.905610
ARIMA(1,0,0)(0,0,2)[7]	Inf
ARIMA(0,0,1)(0,0,2)[7]	1439.537620
ARIMA(0,0,0)(1,0,2)[7]	Inf


--- Re-estimating best models without approximation ---

ARIMA(1,0,0)(0,0,0)[7]+c      -1700.178926

Na etapa de diagnóstico, verificamos se os resíduos são ruído branco e aproximadamente normalmente distribuídos, o que indica que o modelo foi bem especificado. No teste de Ljung-Box, não há evidência suficiente para rejeitar a hipótese nula de autocorrelação dos resíduos. No mesmo sentido, o histograma também aponta para o diagnóstico positivo do modelo.

```
# teste de Ljung-Box
data_fit |>
  dplyr::select(ar1) |>
  gg_tsresiduals()
```





```
augment(data_fit) |>
  dplyr::filter(.model == "ar1") |>
  features(.innov, lbjung_box, lag = 12, dof = 2)
```

```
# A tibble: 1 x 3
  .model lb_stat lb_pvalue
  <chr>   <dbl>   <dbl>
1 ar1     10.9     0.368
```

Conclui-se pela seleção do AR(1), de seguinte equação:

$$y_t = \phi_0 + \phi_1 y_{t-1}$$

## A MODELAGEM BOX-JENKINS: SÉRIE II

Visualizando a série, nota-se imediatamente que não é estacionária por conta da tendência.

```
# importando dados
data = data.frame(
  value = conjunto1[, 6],
```

```


index = yearmonth(
  seq(
    as.Date("2000-01-01"),
    by = "month",
    length.out = length(conjunto1[, 1])
  )
)
) |> tsibble(index = index)

```

```

# plot série
autoplot(data, .vars = value)

```



A estratégia para correção é a diferenciação. Com o teste KPSS, a primeira sugestão é de que a série é estacionária em primeiras diferenças, sem raiz unitária sazonal.

```

# KPSS test e sazonal
data |>
  features(value, unitroot_ndiffs)

```

```

# A tibble: 1 x 1
ndiffs

```

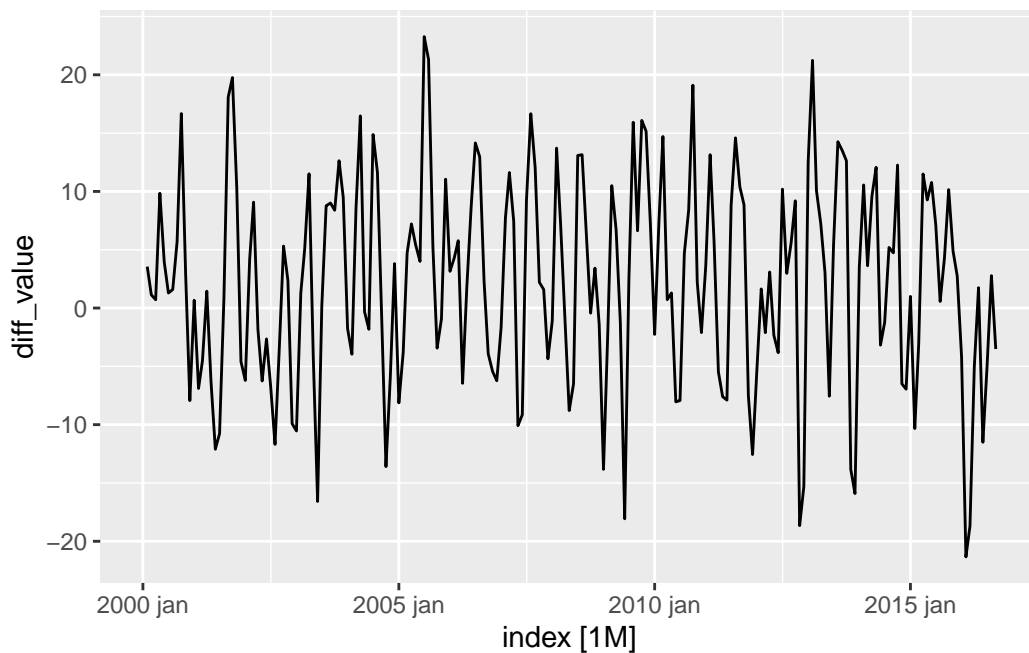
```
<int>  
1      1
```

```
data |>  
  features(value, unitroot_nsdiffs)
```

```
# A tibble: 1 x 1  
  nsdiffs  
    <int>  
1        0
```

```
# diferenciando  
data$diff_value = difference(data$value)  
  
data |>  
  autoplot(.vars = diff_value)
```

Warning: Removed 1 row containing missing values (`geom\_line()`).



Na série diferenciada, não há evidências de raiz unitária tanto nos testes quanto nos gráficos de auto-correlação. Para os testes ADF, iniciei com a especificação com tendência. Não sendo significativo o

coeficiente  $tt$ , passei para a especificação com *drift*, sendo tanto o intercepto quanto  $z.lag.1$  significativos, adoto esta como a correta especificação e, assim como nos testes de Phillips-Perron e KPSS, não há indicativo de raiz unitária.

```
# KPSS test
data |>
  features(diff_value, unitroot_kpss)
```

```
# A tibble: 1 x 2
  kpss_stat kpss_pvalue
    <dbl>     <dbl>
1    0.100      0.1
```

```
# Phillips-Perron test
data |>
  features(diff_value, unitroot_pp)
```

```
# A tibble: 1 x 2
  pp_stat pp_pvalue
    <dbl>     <dbl>
1  -8.31      0.01
```

```
# Augmented-Dickey-Fuller test
data |>
  na.omit() |>
  (\(x) ur.df(x$diff_value, selectlags = "AIC", type = "trend", lags = 12))() |>
  summary()
```

```
#####
# Augmented Dickey-Fuller Test Unit Root Test #
#####
```

Test regression trend

Call:  
lm(formula = z.diff ~ z.lag.1 + 1 + tt + z.diff.lag)

Residuals:

Min	1Q	Median	3Q	Max
-----	----	--------	----	-----

-17.9098 -3.4490 0.5295 4.2145 12.6605

Coefficients:

```
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  1.8911956  1.0526862   1.797 0.074163 .
z.lag.1      -0.7920838  0.2148502  -3.687 0.000304 ***
tt           -0.0007629  0.0084378  -0.090 0.928060
z.diff.lag1   0.5884153  0.2093105   2.811 0.005509 **
z.diff.lag2  -0.1205004  0.2031784  -0.593 0.553909
z.diff.lag3   0.4398270  0.1868001   2.355 0.019674 *
z.diff.lag4  -0.1387117  0.1797101  -0.772 0.441254
z.diff.lag5   0.4108668  0.1660764   2.474 0.014333 *
z.diff.lag6  -0.0965456  0.1633881  -0.591 0.555365
z.diff.lag7   0.3337387  0.1442319   2.314 0.021856 *
z.diff.lag8  -0.0726282  0.1409957  -0.515 0.607138
z.diff.lag9   0.1083096  0.1151223   0.941 0.348116
z.diff.lag10 -0.2205951  0.1098470  -2.008 0.046186 *
z.diff.lag11 -0.0415018  0.0804456  -0.516 0.606588
z.diff.lag12  0.1465711  0.0779209   1.881 0.061658 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 6.112 on 172 degrees of freedom

Multiple R-squared: 0.6248, Adjusted R-squared: 0.5942

F-statistic: 20.46 on 14 and 172 DF, p-value: < 2.2e-16

Value of test-statistic is: -3.6867 4.7354 7.0691

Critical values for test statistics:

```
      1pct  5pct 10pct
tau3 -3.99 -3.43 -3.13
phi2  6.22  4.75  4.07
phi3  8.43  6.49  5.47
```

```
data |>
  na.omit() |>
  (\(x) ur.df(x$diff_value, selectlags = "AIC", type = "drift", lags = 12))() |>
  summary()
```

#####

```
# Augmented Dickey-Fuller Test Unit Root Test #
#####
```

Test regression drift

Call:

```
lm(formula = z.diff ~ z.lag.1 + 1 + z.diff.lag)
```

Residuals:

Min	1Q	Median	3Q	Max
-17.9540	-3.4530	0.4861	4.1709	12.5999

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	1.81874	0.68064	2.672	0.008258	**
z.lag.1	-0.79544	0.21100	-3.770	0.000224	***
z.diff.lag1	0.59173	0.20548	2.880	0.004483	**
z.diff.lag2	-0.11726	0.19941	-0.588	0.557291	
z.diff.lag3	0.44286	0.18324	2.417	0.016697	*
z.diff.lag4	-0.13592	0.17652	-0.770	0.442369	
z.diff.lag5	0.41352	0.16300	2.537	0.012070	*
z.diff.lag6	-0.09400	0.16048	-0.586	0.558816	
z.diff.lag7	0.33595	0.14174	2.370	0.018885	*
z.diff.lag8	-0.07066	0.13890	-0.509	0.611619	
z.diff.lag9	0.10976	0.11366	0.966	0.335547	
z.diff.lag10	-0.21934	0.10866	-2.019	0.045066	*
z.diff.lag11	-0.04082	0.07987	-0.511	0.609893	
z.diff.lag12	0.14712	0.07746	1.899	0.059183	.

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.095 on 173 degrees of freedom

Multiple R-squared: 0.6248, Adjusted R-squared: 0.5966

F-statistic: 22.16 on 13 and 173 DF, p-value: < 2.2e-16


Value of test-statistic is: -3.7698 7.14

Critical values for test statistics:


	1pct	5pct	10pct
tau2	-3.46	-2.88	-2.57
phi1	6.52	4.63	3.81

A seguir, pode-se perceber decaimento nas lags sazonais na ACF, indicando a presença de componente sazonal. Além disso, a análise dos correlogramas sugerem modelos de ordem inferior ou igual a 3 — três picos significativos na ACF e dois na PACF. Porém, a identificação de modelos ARMA é inconclusiva a partir da análise do correlograma exclusivamente.

```
# ACF  
data |> ACF(diff_value, lag_max = 60) |> autoplot()
```



```
# ACF  
data |> PACF(diff_value, lag_max = 60) |> autoplot()
```



Para avaliar candidatos a modelos, realizei uma *grid search* impondo as seguintes restrições:

1.  $d = 1$ ,  $D = 0$  e  $d + D + \text{constante} \leq 2$
2.  $p + q + P + Q \leq 5$

```
data |>
  model(
    search = ARIMA(
      value,
      ic = "aic",
      stepwise = FALSE,
      trace = TRUE,
      order_constraint = (p + q + P + Q <= 5) &
        d == 1 &
        D == 0 &
        constant == 1
    )
  )
```

```
ARIMA(0,1,0)(0,0,0)[12]+c 1437.746540
ARIMA(1,1,0)(0,0,0)[12]+c 1400.743577
ARIMA(2,1,0)(0,0,0)[12]+c 1357.908576
```



ARIMA(3,1,0)(0,0,0)[12]+c	1358.280059
ARIMA(4,1,0)(0,0,0)[12]+c	1355.675578
ARIMA(5,1,0)(0,0,0)[12]+c	1350.559786
ARIMA(0,1,1)(0,0,0)[12]+c	1336.639861
ARIMA(1,1,1)(0,0,0)[12]+c	1340.055957
ARIMA(2,1,1)(0,0,0)[12]+c	1340.606033
ARIMA(3,1,1)(0,0,0)[12]+c	1339.452417
ARIMA(4,1,1)(0,0,0)[12]+c	1337.362051
ARIMA(0,1,2)(0,0,0)[12]+c	1338.490236
ARIMA(1,1,2)(0,0,0)[12]+c	1336.992350
ARIMA(2,1,2)(0,0,0)[12]+c	1338.867528
ARIMA(3,1,2)(0,0,0)[12]+c	1340.520759
ARIMA(0,1,3)(0,0,0)[12]+c	1336.234907
ARIMA(1,1,3)(0,0,0)[12]+c	1337.488084
ARIMA(2,1,3)(0,0,0)[12]+c	Inf
ARIMA(0,1,4)(0,0,0)[12]+c	1336.347799
ARIMA(1,1,4)(0,0,0)[12]+c	1339.128960
ARIMA(0,1,5)(0,0,0)[12]+c	1334.782453
ARIMA(0,1,0)(1,0,0)[12]+c	1402.570657
ARIMA(1,1,0)(1,0,0)[12]+c	1349.906800
ARIMA(2,1,0)(1,0,0)[12]+c	1325.197821
ARIMA(3,1,0)(1,0,0)[12]+c	1321.281273
ARIMA(4,1,0)(1,0,0)[12]+c	1316.627967
ARIMA(0,1,1)(1,0,0)[12]+c	1292.725358
ARIMA(1,1,1)(1,0,0)[12]+c	1300.378066
ARIMA(2,1,1)(1,0,0)[12]+c	1296.747908
ARIMA(3,1,1)(1,0,0)[12]+c	1299.490618
ARIMA(0,1,2)(1,0,0)[12]+c	1294.719267
ARIMA(1,1,2)(1,0,0)[12]+c	1293.050580
ARIMA(2,1,2)(1,0,0)[12]+c	1298.047996
ARIMA(0,1,3)(1,0,0)[12]+c	1296.664829
ARIMA(1,1,3)(1,0,0)[12]+c	1294.284391
ARIMA(0,1,4)(1,0,0)[12]+c	1298.276716
ARIMA(0,1,0)(2,0,0)[12]+c	1397.176416
ARIMA(1,1,0)(2,0,0)[12]+c	1348.459565
ARIMA(2,1,0)(2,0,0)[12]+c	1323.790985
ARIMA(3,1,0)(2,0,0)[12]+c	1318.819544
ARIMA(0,1,1)(2,0,0)[12]+c	1292.680208
ARIMA(1,1,1)(2,0,0)[12]+c	1301.825113
ARIMA(2,1,1)(2,0,0)[12]+c	1296.302358
ARIMA(0,1,2)(2,0,0)[12]+c	1294.562441
ARIMA(1,1,2)(2,0,0)[12]+c	1300.031592
ARIMA(0,1,3)(2,0,0)[12]+c	1296.534677

ARIMA(0,1,0)(0,0,1)[12]+c	1409.998876
ARIMA(1,1,0)(0,0,1)[12]+c	1363.609338
ARIMA(2,1,0)(0,0,1)[12]+c	1334.956399
ARIMA(3,1,0)(0,0,1)[12]+c	1331.560493
ARIMA(4,1,0)(0,0,1)[12]+c	1327.464316
ARIMA(0,1,1)(0,0,1)[12]+c	1301.724023
ARIMA(1,1,1)(0,0,1)[12]+c	1305.101113
ARIMA(2,1,1)(0,0,1)[12]+c	1308.088883
ARIMA(3,1,1)(0,0,1)[12]+c	1309.653416
ARIMA(0,1,2)(0,0,1)[12]+c	1303.721696
ARIMA(1,1,2)(0,0,1)[12]+c	Inf
ARIMA(2,1,2)(0,0,1)[12]+c	Inf
ARIMA(0,1,3)(0,0,1)[12]+c	1304.958754
ARIMA(1,1,3)(0,0,1)[12]+c	Inf
ARIMA(0,1,4)(0,0,1)[12]+c	1306.870600
ARIMA(0,1,0)(1,0,1)[12]+c	1403.448311
ARIMA(1,1,0)(1,0,1)[12]+c	1351.702381
ARIMA(2,1,0)(1,0,1)[12]+c	1327.078436
ARIMA(3,1,0)(1,0,1)[12]+c	1323.279364
ARIMA(0,1,1)(1,0,1)[12]+c	1294.639355
ARIMA(1,1,1)(1,0,1)[12]+c	1302.358114
ARIMA(2,1,1)(1,0,1)[12]+c	1298.747827
ARIMA(0,1,2)(1,0,1)[12]+c	1296.625418
ARIMA(1,1,2)(1,0,1)[12]+c	1295.049111
ARIMA(0,1,3)(1,0,1)[12]+c	1298.576122
ARIMA(0,1,0)(2,0,1)[12]+c	1390.775075
ARIMA(1,1,0)(2,0,1)[12]+c	1348.261266
ARIMA(2,1,0)(2,0,1)[12]+c	1324.585829
ARIMA(0,1,1)(2,0,1)[12]+c	1293.364484
ARIMA(1,1,1)(2,0,1)[12]+c	1298.988623
ARIMA(0,1,2)(2,0,1)[12]+c	1295.002440
ARIMA(0,1,0)(0,0,2)[12]+c	1390.940068
ARIMA(1,1,0)(0,0,2)[12]+c	1347.710071
ARIMA(2,1,0)(0,0,2)[12]+c	1324.551274
ARIMA(3,1,0)(0,0,2)[12]+c	1322.586718
ARIMA(0,1,1)(0,0,2)[12]+c	1291.857638
ARIMA(1,1,1)(0,0,2)[12]+c	1294.617182
ARIMA(2,1,1)(0,0,2)[12]+c	1298.023506
ARIMA(0,1,2)(0,0,2)[12]+c	1293.478278
ARIMA(1,1,2)(0,0,2)[12]+c	Inf
ARIMA(0,1,3)(0,0,2)[12]+c	1295.340501
ARIMA(0,1,0)(1,0,2)[12]+c	1400.796589
ARIMA(1,1,0)(1,0,2)[12]+c	1351.778712

```

ARIMA(2,1,0)(1,0,2)[12]+c 1323.978327
ARIMA(0,1,1)(1,0,2)[12]+c 1296.346899
ARIMA(1,1,1)(1,0,2)[12]+c 1302.532166
ARIMA(0,1,2)(1,0,2)[12]+c 1298.302977
ARIMA(0,1,0)(2,0,2)[12]+c 1391.092422
ARIMA(1,1,0)(2,0,2)[12]+c 1349.476212
ARIMA(0,1,1)(2,0,2)[12]+c 1295.167957

```

--- Re-estimating best models without approximation ---

```

ARIMA(0,1,1)(0,0,2)[12]+c 1298.688664

```

# A mable: 1 x 1

```

      search
      <model>
1 <ARIMA(0,1,1)(0,0,2)[12] w/ drift>

```

A estratégia adotada foi selecionar os modelos com  $\Delta AIC < 2$  em relação ao modelo de menor AIC, além de sua combinação por média simples. São eles:

```

tibble::tribble(
  ~Modelo, ~AIC, ~delta_AIC,
  "ARIMA(0,1,1)(0,0,2)[12]+c", 1291.86, 0,
  "ARIMA(0,1,1)(2,0,0)[12]+c", 1292.68, 0.82,
  "ARIMA(0,1,1)(1,0,0)[12]+c", 1292.73, 0.87,
  "ARIMA(1,1,2)(1,0,0)[12]+c", 1293.05, 1.19,
  "ARIMA(0,1,1)(2,0,1)[12]+c", 1293.36, 1.51,
  "ARIMA(0,1,2)(0,0,2)[12]+c", 1293.48, 1.62
)

```

# A tibble: 6 x 3

	Modelo	AIC	delta_AIC
	<chr>	<dbl>	<dbl>
1	ARIMA(0,1,1)(0,0,2)[12]+c	1292.	0
2	ARIMA(0,1,1)(2,0,0)[12]+c	1293.	0.82
3	ARIMA(0,1,1)(1,0,0)[12]+c	1293.	0.87
4	ARIMA(1,1,2)(1,0,0)[12]+c	1293.	1.19
5	ARIMA(0,1,1)(2,0,1)[12]+c	1293.	1.51
6	ARIMA(0,1,2)(0,0,2)[12]+c	1293.	1.62

```

data_fit = data |>
  model(
    arima011002 = ARIMA(
      value ~ 1 + pdq(0, 1, 1) + PDQ(0, 0, 2)
    ),
    arima011200 = ARIMA(
      value ~ 1 + pdq(0, 1, 1) + PDQ(2, 0, 0)
    ),
    arima011100 = ARIMA(
      value ~ 1 + pdq(0, 1, 1) + PDQ(1, 0, 0)
    ),
    arima112100 = ARIMA(
      value ~ 1 + pdq(1, 1, 2) + PDQ(1, 0, 0)
    ),
    arima011201 = ARIMA(
      value ~ 1 + pdq(0, 1, 1) + PDQ(2, 0, 0)
    ),
    arima012002 = ARIMA(
      value ~ 1 + pdq(0, 1, 2) + PDQ(0, 0, 2)
    )
  ) |>
  dplyr::mutate(combinacao = (
    arima011002 + arima011200 + arima011100 + arima112100 + arima011201 + arima012002
  ) / 6)

```

Na etapa de diagnóstico, todos modelos são considerados aptos para previsão, ao não apresentarem evidências para rejeitar as hipóteses nulas dos testes de Ljung-Box de ausência de autocorrelação serial e ARCH-LM de ausência de heteroscedasticidade condicional.

```

modelos = names(data_fit)
names(modelos) = names(data_fit)

# teste de Ljung-Box
lapply(modelos, function(x) {

  augment(data_fit) |>
  dplyr::filter(.model == x) |>
  features(.innov, ljung_box, lag = 24, dof = 5)
})

```

```
$arima011002
```

```
# A tibble: 1 x 3
  .model      lb_stat lb_pvalue
  <chr>      <dbl>   <dbl>
1 arima011002  22.5     0.260
```

```
$arima011200
# A tibble: 1 x 3
  .model      lb_stat lb_pvalue
  <chr>      <dbl>   <dbl>
1 arima011200  22.5     0.260
```

```
$arima011100
# A tibble: 1 x 3
  .model      lb_stat lb_pvalue
  <chr>      <dbl>   <dbl>
1 arima011100  22.3     0.268
```

```
$arima112100
# A tibble: 1 x 3
  .model      lb_stat lb_pvalue
  <chr>      <dbl>   <dbl>
1 arima112100  20.4     0.370
```

```
$arima011201
# A tibble: 1 x 3
  .model      lb_stat lb_pvalue
  <chr>      <dbl>   <dbl>
1 arima011201  22.5     0.260
```

```
$arima012002
# A tibble: 1 x 3
  .model      lb_stat lb_pvalue
  <chr>      <dbl>   <dbl>
1 arima012002  22.3     0.267
```

```
$combinacao
# A tibble: 1 x 3
  .model      lb_stat lb_pvalue
  <chr>      <dbl>   <dbl>
1 combinacao  22.0     0.284
```

```
# teste ARCH-LM
lapply(modelos, function(x) {

  augment(data_fit) |>
  dplyr::filter(.model == x) |>
  features(.innov, stat_arch_lm, lags = 24)

})
```

```
$arima011002
# A tibble: 1 x 2
  .model      stat_arch_lm
  <chr>         <dbl>
1 arima011002      0.123
```

```
$arima011200
# A tibble: 1 x 2
  .model      stat_arch_lm
  <chr>         <dbl>
1 arima011200      0.127
```

```
$arima011100
# A tibble: 1 x 2
  .model      stat_arch_lm
  <chr>         <dbl>
1 arima011100      0.124
```

```
$arima112100
# A tibble: 1 x 2
  .model      stat_arch_lm
  <chr>         <dbl>
1 arima112100      0.123
```

```
$arima011201
# A tibble: 1 x 2
  .model      stat_arch_lm
  <chr>         <dbl>
1 arima011201      0.127
```

```
$arima012002
# A tibble: 1 x 2
  .model      stat_arch_lm
  <chr>         <dbl>
```

```
1 arima012002      0.122
```

```
$combinacao
```

```
# A tibble: 1 x 2
```

```
  .model      stat_arch_lm
```

```
  <chr>      <dbl>
```

```
1 combinacao      0.125
```

Para testar a performance dos modelos, a série será particionada em 80%-20%, com os modelos treinados na primeira parte e treinada na última. Escolhendo, por fim, o modelo pelo critério menor erro absoluto percentual médio, a seleção ficaria com o SARIMA(0,1,1)(0,0,2), que também foi o de menor AIC.

```
# separando amostra treino
data_treino = subset(data, index <= yearmonth("2013 apr"))

# ajustando o treino
data_treino_fit = data_treino |>
  model(
    arima011002 = ARIMA(
      value ~ 1 + pdq(0, 1, 1) + PDQ(0, 0, 2)
    ),
    arima011200 = ARIMA(
      value ~ 1 + pdq(0, 1, 1) + PDQ(2, 0, 0)
    ),
    arima011100 = ARIMA(
      value ~ 1 + pdq(0, 1, 1) + PDQ(1, 0, 0)
    ),
    arima112100 = ARIMA(
      value ~ 1 + pdq(1, 1, 2) + PDQ(1, 0, 0)
    ),
    arima011201 = ARIMA(
      value ~ 1 + pdq(0, 1, 1) + PDQ(2, 0, 0)
    ),
    arima012002 = ARIMA(
      value ~ 1 + pdq(0, 1, 2) + PDQ(0, 0, 2)
    )
  ) |>
  dplyr::mutate(combinacao = (
    arima011002 + arima011200 + arima011100 + arima112100 + arima011201 + arima012002
  ) / 6)
```

```
# realizando previsões para fora do treino
data_treino_fc = data_treino_fit |>
  fabletools::forecast(h = 42)

# plotando
data_treino_fc |>
  autoplot(
    data |> filter_index("2012-01" ~ .),
    level = NULL
  )
```




Figure 1: avaliação de performance

```
# calculando acurácia
accuracy(data_treino_fc, data)
```

```
# A tibble: 7 x 10
  .model      .type      ME  RMSE  MAE      MPE  MAPE  MASE  RMSSE  ACF1
  <chr>      <chr>    <dbl> <dbl> <dbl>    <dbl> <dbl> <dbl> <dbl> <dbl>
1 arima011002 Test    -4.63  26.1  18.0  -0.524  2.00  0.531  0.655  0.860
2 arima011100 Test    -0.185 24.3  18.2  -0.0348 2.03  0.539  0.610  0.851
3 arima011200 Test     1.36  24.0  18.6   0.136  2.06  0.549  0.603  0.849
```



4	arma011201	Test	1.36	24.0	18.6	0.136	2.06	0.549	0.603	0.849
5	arma012002	Test	-3.61	25.9	18.3	-0.410	2.04	0.541	0.650	0.859
6	arma112100	Test	6.79	24.2	20.2	0.736	2.24	0.597	0.607	0.846
7	combinacao	Test	0.183	24.4	18.5	0.00671	2.06	0.547	0.612	0.852

E sua equação:

$$(1 - L)y_t = (1 - \Theta_1 L^{12} - \Theta_2 L^{24})(1 - \theta_1 L)\epsilon_t$$